

CS 70 - Foundations Of Applied Computer Science

Carter Kruse

Reading Assignment 4

CS 70 - Chapter 6 (Least Squares)

Overview: We now turn our attention to linear systems $A\mathbf{x} = \mathbf{b}$ where A is a “tall” $m \times n$ matrix (with $m > n$), \mathbf{x} is an n -vector of unknowns, and \mathbf{b} is an m -vector.

$$\begin{bmatrix} | & & | \\ \mathbf{a}_1 & \cdots & \mathbf{a}_n \\ | & & | \end{bmatrix} \begin{bmatrix} x_1 \\ \vdots \\ x_n \end{bmatrix} = \begin{bmatrix} b_1 \\ \vdots \\ b_n \end{bmatrix}$$

Such systems are said to be *over-determined* or *over-constrained* because they have more equations/constraints (m) than unknowns (n).

The Least Squares Approach: Instead of requiring no error for some constraints while allowing large error in others, it is often better to consider all m constraints and minimize some sort of average error. In other words, we could compromise and try to find an approximate solution to the tall system so that $A\mathbf{x} \approx \mathbf{b}$.

The Residual Vector: How far off would our approximate solution be? We can quantify this by looking at the difference (or error) between the right-hand side and the left-hand side. We call this error the residual and denote it $\Delta = \mathbf{b} - A\mathbf{x}$.

Given Δ , it seems reasonable to try to make it as small as possible. So one mathematically convenient and intuitive notion of average error is the squared norm

$$|\mathbf{b} - A\mathbf{x}|^2 = |\Delta|^2 = \Delta_1^2 + \cdots + \Delta_m^2$$

So, from all possible m -vectors \mathbf{x} , we could seek the one that makes $|\Delta|$ as small as possible. Mathematically, we can write this as

$$\hat{x} = \operatorname{argmin}_x |\mathbf{b} - A\mathbf{x}|^2$$

where we denote this optimal approximate solution \hat{x} .

Geometric View: The Orthogonality Principle: A geometric interpretation in terms of the columns provides a particularly useful starting point: the smallest residual vector must be the one that is perpendicular to the span of A . In other words, Δ must be perpendicular to each column of A .

Algebraic View: Solution Via Calculus: Another way to derive the solution is to treat the residual as a function $E(\mathbf{x}) = |\mathbf{b} - A\mathbf{x}|^2$, and find the minimum of this function by setting its derivatives equal to zero.

General Solution: We turn our attention back to the general system $A\mathbf{x} = \mathbf{b}$ where A is an $m \times n$ matrix (with no restriction on m or n), \mathbf{x} is an n -vector of unknowns, and \mathbf{b} is an m -vector. With n columns and unknowns, the principle of orthogonality we applied above leads to n orthogonality constraints, which we could express in exactly the same way, by requiring that the inner products between the columns of A and the residual are zero:

$$\begin{bmatrix} - & a_1^T & - \\ - & a_2^T & - \\ & \dots & \\ - & a_m^T & - \end{bmatrix} (\mathbf{b} - A\mathbf{x}) = \mathbf{0}$$

This implies $A^T(\mathbf{b} - A\mathbf{x}) = \mathbf{0} \rightarrow A^T A\mathbf{x} = A^T \mathbf{b}$.

Normal Equations: The best least squares solution to the system $A\mathbf{x} \approx \mathbf{b}$, i.e. the \mathbf{x} that minimizes the residual norm $|\mathbf{b} - A\mathbf{x}|$, satisfies

$$A^T A\mathbf{x} = A^T \mathbf{b}$$

$A^T A$ is a matrix - sometimes called the *Gramm matrix* - which is square. It has the same number of rows and columns as the number of unknowns in \mathbf{x} . The normal equations essentially convert tall over-constrained systems into square systems, allowing us to use algorithms for solving square systems to obtain the least squares solution.

Moore-Penrose Inverse: It is possible to prove that as long as the columns of A are linearly independent, then $A^T A$ will be invertible. This means we can express the least squares solution mathematically as

$$\hat{\mathbf{x}} = (A^T A)^{-1} A^T \mathbf{b}$$

This is called the *pseudo-inverse* because it acts much like the inverse of a square matrix, but generalized to a non-square matrix.

See the course notes for further information about the dot product, projection onto a line, and projection onto a subspace.

CS 70 - Chapter 7 (Data Fitting)

Data: We have a set of data samples $(x_1, y_1), (x_2, y_2), \dots, (x_n, y_n)$ in which x_i is a k -vector and y_i is a scalar.

Function: We believe x and y are related. Their relation can be described by a function $f: \mathbb{R}^n \rightarrow \mathbb{R}$ that maps from x to y :

$$y = f(x)$$

Model: We do not know the expression of $f(x)$. So, we aim to approximate $f(x)$ by fitting a model $\hat{f}(x)$ between x and y :

$$\hat{y} = \hat{f}(x)$$

where $\hat{f}: \mathbb{R} \rightarrow \mathbb{R}$. The hat appearing over f is traditional notation that suggests that the function \hat{f} is an approximation of the function f .

Linear Model & Basis: We will focus on a specific form for the model, which has the form

$$\hat{f}(x) = \theta_1 f_1(x) + \theta_2 f_2(x) + \dots + \theta_m f_m(x)$$

where $f_i(x): \mathbb{R}^k \rightarrow \mathbb{R}$ are basis functions or feature mappings that we choose, and θ_i are the model parameters that we choose.

Error: For data sample i , our model predicts the value $\hat{y}_i = \hat{f}(x_i)$, so the prediction error or residual for this data point is

$$\Delta_i = y_i - \hat{y}_i$$

Build Least-Squares Data-Fitting System: We solve a linear system in the least squares sense to find the best parameters $\theta_1, \theta_2, \dots, \theta_n$ for the basis functions f_1, f_2, \dots, f_m given n data samples. The problem can be written in matrix form as $\mathbf{y} = A\boldsymbol{\theta}$.

$$\begin{bmatrix} y_1 \\ y_2 \\ \dots \\ y_n \end{bmatrix} = \begin{bmatrix} f_1(x_1) & f_2(x_1) & \dots & f_m(x_1) \\ f_1(x_2) & f_2(x_2) & \dots & f_m(x_2) \\ \dots & \dots & \dots & \dots \\ f_1(x_n) & f_2(x_n) & \dots & f_m(x_n) \end{bmatrix} \begin{bmatrix} \theta_1 \\ \theta_2 \\ \dots \\ \theta_m \end{bmatrix}$$

This gives us a tall matrix A with its elements calculated by the basis functions taking different data samples x . The best values of $\boldsymbol{\theta}$ in the least squares sense can be calculated as $\boldsymbol{\theta} = (A^T A)^{-1} A^T \mathbf{y}$. The least squares solution minimizes the L_2 norm of the residual $|\Delta| = |\mathbf{y} - \hat{\mathbf{y}}| = |\mathbf{y} - A\boldsymbol{\theta}|$.

CS 70 - Chapter 8 (Gram-Schmidt & QR Decomposition)

Orthonormal Vectors: A collection of vectors $\mathbf{q}_1, \mathbf{q}_2, \dots, \mathbf{q}_n$ are orthonormal if

$$\mathbf{q}_i^T \mathbf{q}_j = \begin{cases} 1 & i = j \\ 0 & i \neq j \end{cases}$$

If we write the collection of vectors as the columns of a matrix, i.e. $Q = [\mathbf{q}_1 \mathbf{q}_2 \dots \mathbf{q}_n]$, then we can write this property succinctly as

$$Q^T Q = I$$

which implies that $Q^{-1} = Q^T$, indicating that we only need to perform a matrix transpose when we want to calculate the inverse of an orthonormal matrix.

Gram-Schmidt Algorithm: Idea: Incrementally build an orthonormal basis one vector at a time. For each newly added vector, we project out the components from the existing orthonormal basis vectors and normalize it after all the projections.

Algorithm: As summarized in the following pseudo-code, you (1) pick one vector each time, (2) project out all the components on the existing orthonormal basis vectors, and (3) normalize.

QR Decomposition: Idea: For a matrix A , conduct the Gram-Schmidt algorithm for the collection of its column vectors $\mathbf{a}_1, \mathbf{a}_2, \dots, \mathbf{a}_n$ to get the orthonormal matrix Q , and at the same time record all the coefficients during the procedure in another upper triangle matrix R .

Note that the QR decomposition can be done by generating each column of both Q and R in sequence, by following the iterative steps in the Gram-Schmidt algorithm. That is, in the first iteration, we calculate the first column of Q (\mathbf{q}_1) and the first column of R ($|\mathbf{q}_1|$). In the second iteration, we calculate the second column of Q (\mathbf{q}_2) and the second column of R ($(\mathbf{a}_2^T \mathbf{q}_1) \mathbf{q}_1$ and $|\hat{\mathbf{q}}_2|$), etc.

QR decomposition can be used to solve least-squares problems. For a least-squares system $A\mathbf{x} = \mathbf{b}$, we first factorize A as $A = QR$, then

- Solve $Q\mathbf{y} = \mathbf{b}$ by $\mathbf{y} = Q^T \mathbf{b}$ (recall that we have $Q^{-1} = Q^T$ for an orthogonal matrix).
- Solve $R\mathbf{x} = \mathbf{y}$ using backward substitution.